



residual encoder: 7 stages, stride 2, 3^3 kernels, instance normalisation decoder: transposed convolution + 1 block per stage

skip connections (concatenate)

CT patch, 1 channel
 $384 \times 256 \times 256$
 $0.80 \times 0.75 \times 0.75$ mm



a posterior, not a verdict

deep supervision: class-weighted Dice + cross-entropy at four scales

ResEnc-L (nnU-Net v2), 100 GB plan on an H200; batch 2; SGD with Nesterov momentum 0.99, polynomial LR from $1e-2$, 500 epochs
 sampler: half of every queue from the 39 variant carriers, patch forced onto L6 or the lumbar rib